

Programming Interview Prep



Institute for Computational and Mathematical Engineering
Stanford University
Andreas Santucci

Contents

1	A Note on Complexity and Leveling	2
2	Intermediate Programming Questions	3
2.1	Balanced String Split	3
2.2	Intersection of Three Arrays	3
2.3	Two-Sum	4
2.4	Counting Complete Tree Nodes	5
2.5	Majority Element	6
3	Intermediate SQL	8
3.1	With Clauses and Sub-Queries	8
3.2	Window Expressions	9
3.3	Applied Problem: Ranking Scores	10
4	More Programming Questions (*for those who took 212 already)	11
4.1	Warmup: defanging an IP address	11
4.2	Remove outermost parentheses	11
4.3	Merge two binary trees	12

1 A Note on Complexity and Leveling

Interviewing for a Fixed Level When new grads enter the job market, they are all considered for a level that is based on their degree, e.g. level L for an M.S. and $L + 1$ for a PhD. It doesn't matter how proficient you are, given your credentials you will be interviewing for a new grad position. This means that the interview questions will be scoped accordingly. Remember that in a 45 minute interview, you can expect a warmup question to be solvable in under 15 minutes and a more challenging question which lends itself to open-ended follow up questions for the remainder of the technical interview. It is not likely that you will have time to ask your interviewer questions at the end of the interview, because they will likely be continually assessing your ability by probing with increasingly difficult questions. This means that it's OK if you don't finish all the interview questions; don't let it hurt your morale and instead focus on the the next interview!

Showing Leadership It is still tough to earn the job offer. How can you showcase your understanding of programming and quantitative reasoning?

- For questions about your research.
 - Focus on describing your work in *simple*, easy to understand language that any person would be able to understand. If you used a complex method, be prepared to explain why simpler methods didn't work.
- Ask questions about the problem statement.
 - Clarify ambiguities that the interviewer may have intentionally left open to interpretation.
 - Determine what constraints you are working under; sometimes the solution can be made simpler if you have permission to use certain methods or tools.
- Demonstrate knowledge of tradeoffs.
 - Each problem likely has multiple solutions. Can you spot them and mention their relative merits during the interview? This can showcase your broad knowledge. It's enough to mention various tradeoffs before proceeding with an an approach that is tied for efficiency.
- Explain what you're thinking, always!
 - One of the biggest mistakes interviewers make is solving problems silently; this comes across as a hard pass in that people don't want to work with peers who don't explain their solutions eloquently.
 - Make sure to explain what you're doing at each step, out loud, and consistently pause to ask the interviewer if they have any questions on why you're taking the approach proposed.

2 Intermediate Programming Questions

2.1 Balanced String Split

Balanced strings are those who have equal quantity of 'L' and 'R' characters.

Given a balanced string *s* split it in the maximum amount of balanced strings.

Return the maximum amount of splitted balanced strings.

```

1  int balancedStringSplit(string s) {
2      int balance {};
3      int num_splits {};
4      for (char c : s) {
5          c == 'L' ? ++balance : --balance;
6          if (balance == 0) num_splits++;
7      }
8      return num_splits;
9  }

```

Idea behind the algorithm The trick is to maintain a 'balance' variable which we will increment by unit value when we run into an 'L' and we will decrement by unit value when we run into an 'R'. Each time the balance var "hits" zero, we know that we can split off a piece of the string into a **balancedSubString**. One *key argument to make* is that the remainder is balanced by virtue of the original input being balanced.

Complexity It's easy to see that our algorithm requires $\mathcal{O}(1)$ storage and $\mathcal{O}(n)$ work.

2.2 Intersection of Three Arrays

Find the intersection of three input arrays. Return the output in sorted order.

```

1  vector<int> arraysIntersection(vector<int>& arr1, vector<int>& arr2, vector<int>& arr3) {
2      map<int, int> counts;
3      for (const auto& val : arr1) counts[val] += 1;
4      for (const auto& val : arr2) counts[val] += 1;
5      for (const auto& val : arr3) counts[val] += 1;
6      vector<int> result;
7      for (const auto& kv : counts)
8          if (kv.second == 3) result.push_back(kv.first);
9      // We don't need to sort, since we're already using an ordered map.
10     return result;
11 }

```

Idea behind the algorithm Create a map which counts how many times we see a particular value *across* all three input arrays. In the end, the values which are associated with a count of three correspond to being in all three input arrays.

Analysis Each insertion costs logarithmic time with respect to the size of the *ordered* map. If n is the size of the largest input vector, then we require $O(n \log n)$ work to populate our ordered map. Populating the output container only requires a linear traversal. Therefore, the total work is given by $O(n \log n)$.

Emphasizing why *ordered* map is necessary The problem statement asked for the output in sorted order. Therefore, it's easy enough to use an *ordered map* which means that when we iterate over it, the values of our original input(s) (which are now keys in the map) will be in the appropriate order in our last step. If we insisted on using an *unordered map*, we wouldn't gain anything, since although creating the counts would only require linear time, we'd need to sort our output which would bring us back up to $O(n \log n)$ work.

Follow Up It's trivial to find the intersection of k arrays using the same technique; use a `std::vector<std::map<int,int>>`.

2.3 Two-Sum

Given an array of integers, return indices of the two numbers such that they add up to a specific target. You may assume that each input would have exactly one solution, and you may not use the same element twice. Don't worry about possible overflow.

(Brute-Force) Solution It's almost always worth it to at least mention a brute force solution before proceeding to an efficient implementation.

```

1 vector<int> twoSum(vector<int>& nums, int target) {
2     for (unsigned i = 0; i < nums.length() - 1; ++i)
3         for (unsigned j = i + 1; j < nums.length(); ++j)
4             if (nums[i] + nums[j] == target)
5                 return {i, j};
6     return {-1, -1};
7 }
```

The brute force solution requires $O(n^2)$ work. It's also possible to come up with a **linear time algorithm**, i.e. we can find the indices in $O(n)$ time with $O(n)$ storage.

```

1 vector<int> twoSum(vector<int>& nums, int target) {
2     unordered_map<int, int> m; // Associate values in 'nums' with their corresponding indices.
3     m[nums[0]] = 0;
4     for (int i = 1; i < nums.size(); ++i) {
5         // find is  $O(1)$  average complexity for unordered map.
6         // Search contains a first (key) and second (val) element.
7         auto search = m.find(target - nums[i]);
8         if (search != m.end())
9             return {search->second, i};
10        m[nums[i]] = i;
11    }
12    return {-1, -1};
13 }
```

Idea behind the linear time algorithm Addition is commutative, i.e. $a + b = b + a$. Therefore, it doesn't matter what order we consider our inputs. For each element in the input, we can compute `target` $- x$ and see if this candidate value is contained in our container. If it is, it means we've already come across an element which, when added to the current element x that we're iterating over, adds to our `target` value.

Complexity analysis We make a linear traversal over the input. For each element, we perform a subtraction in constant time and then use a `find` operation on an *unordered* map to determine if the two-sum achieves the target value in constant time. The `if`-statement and the final insertion into the unordered map are both $\mathcal{O}(1)$ as well, and so the entire algorithm is $\mathcal{O}(n)$.

Extension: three sum A natural follow-up question is to come up with an algorithm for finding a three-sum that achieves a target value. This can be an exercise for the reader.

2.4 Counting Complete Tree Nodes

There are a myriad of tree traversal algorithms that are fun to cover in interviews. We present a *very* simple one here.

Problem Given a complete binary tree, count the number of nodes. In a complete binary tree every level, except possibly the last, is completely filled, and all nodes in the last level are as far left as possible. It can have between 1 and 2^h nodes inclusive at the last level h .

Example

Input:

```

      1
     / \
    2   3
   / \ /
  4  5 6

```

Correct output: 6.

```

/**
 * Definition for a binary tree node.
 * struct TreeNode {
 *     int val;
 *     TreeNode *left;

```

```

*    TreeNode *right;
*    TreeNode(int x) : val(x), left(NULL), right(NULL) {}
* };
*/
int countNodes(TreeNode* root) {
    if (!root) return 0;
    return 1 + countNodes(root->left) + countNodes(root->right);
}

```

Complexity Analysis Our algorithm touches each node once and performs a simple arithmetic computation with each invocation. We therefore have a linear traversal (with respect to the number of nodes in the input graph).

2.5 Majority Element

Given an array of size n , find the majority element. The majority element is the element that appears more than $\lfloor n/2 \rfloor$ times. You may assume that the array is non-empty and the majority element always exists in the array.

Example 1:

Input: [3,2,3]

Output: 3

Example 2:

Input: [2,2,1,1,1,2,2]

Output: 2

Sorting as a means to an end One simple approach we can take is to use a sort as a means to an end. Since the majority element appears *more* than $\lfloor \frac{n}{2} \rfloor$ times, after sorting we can simply inspect the middle element.

```

1 int majorityElement(vector<int>& nums) {
2     sort(nums.begin(), nums.end());
3     return nums[nums.size() / 2];
4 }

```

For any reasonable `sort` implementation we can say this is an $\mathcal{O}(n \log n)$ algorithm.

Divide and Conquer Next we explore a divide and conquer approach; the context is that your interviewer doesn't want you to use a built-in sorting routine, but instead wants to see you do some of the work at a lower level.

Hint: suppose we are told the majority element for the left half and the right half of the array. If these values coincide, we've found our majority element and we're done. Otherwise,

iterative over the sub-input in linear time and see which of the two occurs more often, that's the majority element. The base case is when there's only one element in the array, then trivially it itself is the majority element.

Intuition for why linear scan is necessary To see why we need the linear traversal, think about an input where an element occurs all but one time in the left half of the array, and the right half consists of it appearing only a couple times (with some other element appearing a "majority" number of times) then we clearly need to figure out which of the two candidate majority elements is the right one.

```
1  int majorityElement(vector<int>& nums) { return majorityElement(nums, 0, nums.size() - 1); }
```

We then define a recursive work-horse.

```
1  int majorityElement(const vector<int>& nums, int l, int r) {
2      if (l == r)
3          return nums[l];
4      int mid = (r-l)/2 + 1;
5      int l_max = majorityElement(nums, l, mid);
6      int r_max = majorityElement(nums, mid+1, r);
7      if (l_max == r_max) // If they agree, we're done.
8          return l_max;
9      int l_count{}, r_count{}; // Else count which occurs more often.
10     for (int k = l; k <= r; k++) { // Iterate over the _sub_-array.
11         if (nums[k] == l_max) l_count++;
12         if (nums[k] == r_max) r_count++;
13     }
14     return l_count > r_count ? l_max : r_max;
15 }
```

Complexity analysis Recurrence relation for work required:

$$T(N) = 2 * T(N/2) + \mathcal{O}(N) = \mathcal{O}(N \log N).$$

You can use [Master's theorem](#) or unroll the recurrence relation by hand to arrive at the asymptotic complexity.

3 Intermediate SQL

3.1 With Clauses and Sub-Queries

A **with** clause lets you name a *sub-query*. I.e. it's possible to write a SQL program with more than one **SELECT ... FROM ... WHERE ...**. Suppose you have the following two tables, the first maps customer id's to transaction amounts, and it's called **transactions**:

id	amt
901	8.00
901	9.00
902	1.00
902	2.00

and the second maps customer id's to names, called **customer_info**:

id	name
901	Andre
902	Nick

What if we want to calculate the average transaction amount by customer, and then display the result alongside their name? We can use a **with** clause.

```
WITH avg_amt AS (  
  SELECT  
    id  
    , AVG(amt) AS avg_transaction_price  
  FROM  
    transactions  
  GROUP BY 1  
) SELECT  
  name  
  , avg_transaction_price AS amt  
FROM  
  avg_amt JOIN customer_info ON avg_amt.id = customer_info.id;
```

The result would look like:

name	amt
Andre	8.50
Nick	1.50

3.2 Window Expressions

An aggregation function takes in a sequence of observations and collapses it into a single number. What if we want to take in a sequence of values and output another sequence of values, e.g. a cumulative sum? We can use a window function for this.

Suppose we have a race tournament where there are multiple stages of racing, and we'd like to calculate (ex-post) each participant's final ranking at each stage of the race, based on all previous rounds. The easy way to do this would be to generate a column that describes the cumulative time spent up until the given stage. Suppose our data looks like `races`:

id	trial	time
A	0	0:30
A	1	0:10
A	2	0:05
...	...	
B	0	0:28
B	1	0:07
B	2	0:05

We already know how we could calculate the *total* time spent across all races for a given participant-id using a group-by clause, e.g. we could rank the participants by their total times via `SELECT id, SUM(time) AS final_time FROM races GROUP BY 1 ORDER BY 2;`. But what about the cumulative sum of times? We need a function that takes in a window of n observations and returns another sequence of n observations, rather than a typical aggregation method which collapses n observations into a scalar value...

```
SELECT
  id
  , trial
  , SUM(time) OVER(PARTITION BY id ORDER BY trial) AS cum_time
FROM
  races;
```

This would return something like:

id	trial	cum_time
A	0	0:30
A	1	0:40
A	2	0:45
...
B	0	0:28
B	1	0:35
B	2	0:40

3.3 Applied Problem: Ranking Scores

Write a SQL query to rank scores. If there is a tie between two scores, both should have the same ranking. Note that after a tie, the next ranking number should be the next consecutive integer value. In other words, there should be no "holes" between ranks. Given a **Scores** table:

+-----+-----+			+-----+-----+		
Id	Score		score	Rank	
+-----+-----+			+-----+-----+		
1	3.50		4.00	1	
2	3.65		4.00	1	
3	4.00	-->	3.85	2	
4	3.85		3.65	3	
5	4.00		3.65	3	
6	3.65		3.50	4	
+-----+-----+			+-----+-----+		

Solution The problem can be solved by using both sub-queries and a window clause.

```
WITH distinct_scores AS (
  SELECT DISTINCT
    Score AS score
  FROM
    Scores
), ranked_scores AS (
  SELECT
    score
    , ROW_NUMBER() OVER(ORDER BY score DESC) AS 'rank'
  FROM distinct_scores
)
SELECT
  score
  , 'rank'
FROM
  Scores JOIN ranked_scores USING(score)
ORDER BY 1 DESC;
```

Explanation of query We first query for the distinct scores, and in this way we can later get a rank ordering without having to worry about duplicate scores getting assigned different ranks. In the next step, we actually rank our scores in descending order by utilizing a window function. Lastly, we join our original scores with our crafted ranks and order the result.

4 More Programming Questions (*for those who took 212 already)

4.1 Warmup: defanging an IP address

Given a valid (IPv4) IP address, return a “defanged” version of that IP address. A defanged IP address replaces every period “.” with “[.]”.

```

1  string defangIPAddr(string address) {
2      int idx = address.find('.');
3      while (idx != string::npos) {
4          // The tricky part is to realize that insertion
5          // happens exactly AT the index specified.
6          address.insert(idx, 1, '[');
7          // Further, this means that after insertion,
8          // the '.' character is found at position
9          // idx+1, so to insert a ']' after the period
10         // we need to insert at posn idx + 2.
11         address.insert(idx + 2, 1, ']');
12         // By similar reasoning, our '.' character
13         // got shifted into position idx+1, so start
14         // our next search at position idx+2.
15         idx = address.find('.', idx+2);
16     }
17     return address;

```

Reasoning about the solution You can expect to see a basic warmup question like this that merely tests your ability to write a *very* simple program. The only tricky part of this solution is that it requires some diligent tracking of indices as we perform insertions, since we’re doing the routine “in-place”. Since calling *insert* requires $\mathcal{O}(n)$ work, we have a solution that is potentially quadratic in runtime. We could bring this down further but this is just a warmup question.

4.2 Remove outermost parentheses

A valid parentheses string is either empty “”, “‘(’ + A + ‘)’”, or $A + B$ where A and B are valid parentheses strings, and $+$ represents string concatenation. For example, the empty string, “()”, “((()))”, and “(()())” are all valid parentheses strings.

A valid parentheses string is *primitive* if it is non-empty, *and* there doesn’t exist a way to split it into $S = A + B$, with A and B nonempty valid parentheses strings.

Given a valid parentheses string S , consider its primitive decomposition $S = P_1 + P_2 + \dots + P_k$, where P_i are primitive valid parentheses strings.

Return S after removing the outermost parentheses of every primitive string in the primitive decomposition of S .

```

1  string removeOuterParentheses(string S) {
2      // We do a linear traversal and maintain a balance counter for
3      // the parentheses, realizing that in the primitive decomposition

```

```

4  // the number of '(' and ')' must match.
5  // Each time the balance goes to zero, we've found a primitive substring
6  // and we can append everything but the outer parentheses to the result.
7  int beg{}, idx{}, balance{};
8  string result;
9  while (idx < S.size()) {
10     if (S[idx] == '(') balance++;
11     else if (S[idx] == ')') balance--;
12     if (!balance) {
13         for (int i = beg+1; i < idx; ++i)
14             result.append(1, S[i]);
15         beg = idx+1;
16     }
17     idx++;
18 }
19 return result;
20 }

```

Complexity Analysis This should remind you of the question on “Balanced String Split”, it’s just a bit more abstract. We do a linear scan, taking care to touch each character at most once. Our time and space complexity are both linear.

4.3 Merge two binary trees

Given a two binary trees `root1` and `root2`, you need to merge the two together as follows:

- In your head, imagine overlaying the two trees on top of each other, by conceptually hanging them from the root node from the same fixed point.
- If two nodes overlap, take the output node value to be the sum of the two input nodes.
- Otherwise, the non-null node is used as the node of the new tree.

```

1  /**
2   * Definition for a binary tree node.
3   * struct TreeNode {
4   *     int val;
5   *     TreeNode *left;
6   *     TreeNode *right;
7   *     TreeNode() : val(0), left(nullptr), right(nullptr) {}
8   *     TreeNode(int x) : val(x), left(nullptr), right(nullptr) {}
9   *     TreeNode(int x, TreeNode *left, TreeNode *right) : val(x), left(left), right(right) {}
10  * };
11  */
12  TreeNode* mergeTrees(TreeNode* root1, TreeNode* root2) {
13      // We will maintain that t1 is the merged tree. So, if it's NULL
14      // we'd prefer to swap with the other tree (potentially NULL as well, but that's OK).
15      if (!t1) swap(t1, t2);
16      if (!t2) return t1; // If only one tree left, return it.
17      // Otherwise, both t1 and t2 are non-NULL, and we simply merge the values into t1.
18      t1->val += t2->val;
19      // We then recurse, setting left and right children to the result of a call to
20      // mergeTrees.
21      t1->left = mergeTrees(t1->left, t2->left);
22      t1->right = mergeTrees(t1->right, t2->right);
23      return t1;
24  }

```

Complexity analysis We touch each node once: the time complexity of our algorithm is linear in the size of tree nodes. Because we use recursion, the depth of recursion and consequently the stack-frames required to store the (intermediate) results also grow linearly with respect to the number of nodes in the input tree.

Problem Appendix

If you're interested in covering more interview problems, check out CME 212 if you haven't already; we cover at least one interview programming problem each lecture! Here are a few interesting problems to solve otherwise.

- [Two pointer algorithms](#)
 - [Merge two sorted arrays, retain sorted order \(*in linear time\)](#)
 - [Minimum size subarray sum](#)
- [Recursion](#)
 - [Last stone weight](#)
- [Sliding window technique](#)
 - [Find length-k substrings with no repeated characters](#)
 - [Longest substring without repeating characters](#)

References

- [1] [LeetCode](#) Accessed Fall 2020